

# K-Means Clustering

In this exercise we will use KMeans clustering to create clusters in the iris dataset and then compare them with the original to see how accurate they are.

## Importing Necessary Libraries

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

executed in 2.16s, finished 18:14:02 2023-04-19

## Data Preparation

In [2]:

```
iris = pd.read_csv('iris.csv')
iris.head()
```

executed in 34ms, finished 18:14:02 2023-04-19

Out[2]:

	sepalength	sepalwidth	petallength	petalwidth	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

In [3]:

iris.info()

executed in 22ms, finished 18:14:02 2023-04-19

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepallength      150 non-null    float64
1   sepalwidth       150 non-null    float64
2   petallength      150 non-null    float64
3   petalwidth       150 non-null    float64
4   class            150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

In [4]:

df = iris.copy()

executed in 11ms, finished 18:14:02 2023-04-19

We copy the iris dataset to perform our operations on it and compare it with the original dataset.

In [5]:

df['class'].unique()

executed in 10ms, finished 18:14:02 2023-04-19

Out[5]:

```
array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

We see that there are 3 unique species. Now we will check whether our clustering model predicts the same.

In [6]:

```
df.drop(['class'], axis = 1, inplace = True)
df.head()
```

executed in 31ms, finished 18:14:02 2023-04-19

Out[6]:

	sepallength	sepalwidth	petallength	petalwidth
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

We dropped the `class` column as it was redundant for our exercise.

## Model Development

First we will run the Hopkin's Clustering Test on the dataset. We begin by defining the `hopkins()` function.

In [7]:

```
from sklearn.neighbors import NearestNeighbors
from random import sample
from numpy.random import uniform
from math import isnan

def hopkins(X):
    d = X.shape[1] # columns
    n = len(X) # rows
    m = int(0.1 * n)
    nbrs = NearestNeighbors(n_neighbors=1).fit(X.values)

    rand_X = sample(range(0, n, 1), m)

    ujd = []
    wjd = []
    for j in range(0, m):
        u_dist, _ = nbrs.kneighbors(uniform(np.amin(X,axis=0),np.amax(X,axis=0),d).reshape(
        ujd.append(u_dist[0][1])
        w_dist, _ = nbrs.kneighbors(X.iloc[rand_X[j]].values.reshape(1, -1), 2, return_dist
        wjd.append(w_dist[0][1])

    H = sum(ujd) / (sum(ujd) + sum(wjd))
    if isnan(H):
        print(ujd, wjd)
        H = 0

    return H
```

executed in 210ms, finished 18:14:02 2023-04-19

Since Hopkin's test produces a slightly different value each time we will iterate it for 1000 times and then take its mean.

In [8]:

```
hopkins_stat = []
for i in range(1000):
    hopkins_stat.append(hopkins(df))

sum(hopkins_stat)/len(hopkins_stat)
```

executed in 50.7s, finished 18:14:53 2023-04-19

Out[8]:

0.842680849830881

The mean value is 0.84 which means our dataset has a high tendency for clustering. Now we will train the model. We will use elbow curve method for determining the number of clusters.

In [9]:

```
from sklearn.cluster import KMeans
```

executed in 125ms, finished 18:14:53 2023-04-19

If we set number of clusters as 1 we will get an error so we will calculate it manually. For `n_clusters = 1` the data will have a single centroid and we have to calculate the sum of the square of distance between the centroid and the remaining data points. First we will calculate the centroid.

In [10]:

```
centroid = np.mean(df, axis = 0)
centroid = centroid.values
centroid
```

executed in 28ms, finished 18:14:53 2023-04-19

Out[10]:

```
array([5.84333333, 3.054      , 3.75866667, 1.19866667])
```

Now we will create two UDFs to calculate the distance between 2 points and the sum of those distances.

In [11]:

```
def euclidean_distance(point1, point2):
    """Calculate the Euclidean distance between two points."""
    return np.sqrt(np.sum((point1 - point2)**2))

def sum_euclidean_distance(centroid, data_points):
    """Calculate the sum of Euclidean distances between the centroid and all data points."""
    sum_distance = 0
    for point in data_points:
        sum_distance += euclidean_distance(centroid, point)
    return sum_distance
```

executed in 24ms, finished 18:14:53 2023-04-19

In [12]:

```
sum_euclidean_distance(centroid, df.values)
```

executed in 34ms, finished 18:14:53 2023-04-19

Out[12]:

```
291.4551238555539
```

The inertia for `n_clusters = 1` is 291.45 . We will insert that value in the list of inertias.

In [13]:

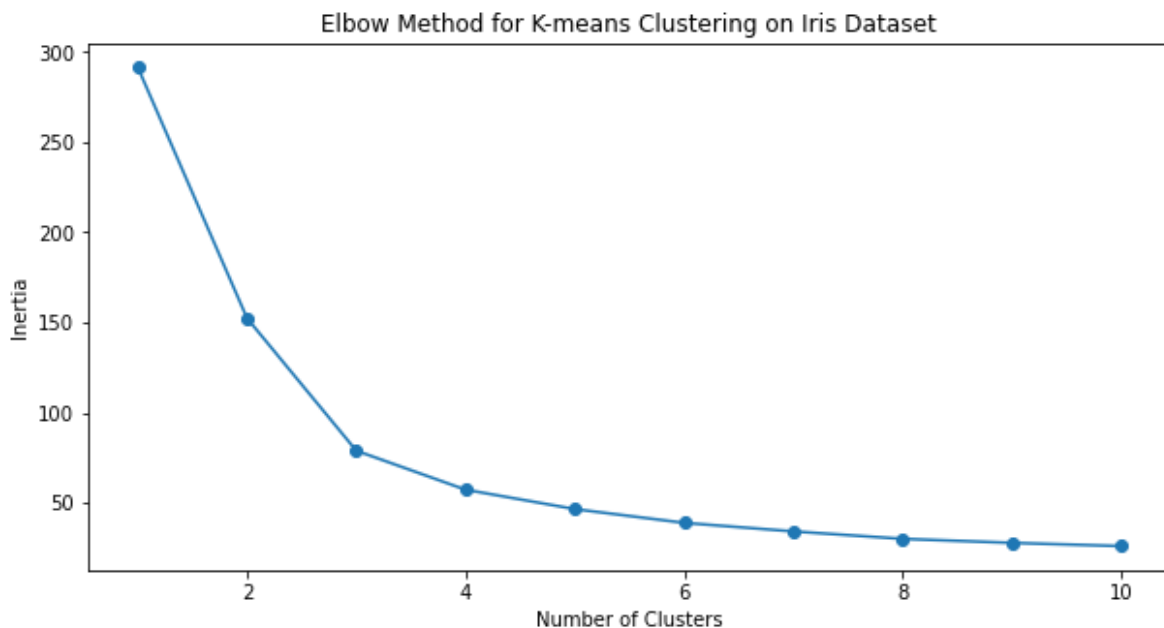
```
# Define a list of number of clusters to try
n_clusters = range(2, 11)

# Initialize an list with inertia value for n = 1 to store the inertia values for other n
inertias = [291.45]

# Perform K-means clustering for different number of clusters
for n in n_clusters:
    kmeans = KMeans(n_clusters=n, init='k-means++', max_iter=300, random_state=0)
    kmeans.fit(df)
    inertias.append(kmeans.inertia_)

# Plot the inertia values vs. number of clusters
plt.figure(figsize = (10,5))
plt.plot(range(1,11), inertias, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method for K-means Clustering on Iris Dataset')
plt.show()
```

executed in 1.25s, finished 18:14:54 2023-04-19



The elbow point is at  $n = 3$  which is in accordance with our dataset. We will now train the model for 3 clusters and then compare the results with original dataset.

In [14]:

```
kmeans_3 = KMeans(n_clusters = 3, init='k-means++', max_iter=300, random_state=0)
kmeans_3.fit(df)
```

executed in 113ms, finished 18:14:54 2023-04-19

Out[14]:

```
KMeans(n_clusters=3, random_state=0)
```



In [18]:


```
plt.figure(figsize=(30,8))

plt.subplot(1,2,1)
sns.scatterplot(x = 'sepalength', y = 'sepalwidth', data = iris, hue = 'class', palette =
plt.title('Iris Dataset', fontsize = 14)
plt.legend(title = 'class', bbox_to_anchor=(1, 1), loc = 0)

plt.subplot(1,2,2)
sns.scatterplot(x = 'sepalength', y = 'sepalwidth', data = df, hue = 'cluster', palette =
plt.title('KMeans Model', fontsize = 14)
plt.legend(title = 'cluster', bbox_to_anchor=(1, 1), loc = 0)

plt.suptitle("Sepallength vs Sepalwidth", fontsize = 20)
plt.show()
```

executed in 1.29s, finished 18:14:56 2023-04-19



Our model did a great job at creating the clusters. Now we will check the petallength vs petalwidth graph.

In [19]:


```
plt.figure(figsize=(30,8))

plt.subplot(1,2,1)
sns.scatterplot(x = 'petallength', y = 'petalwidth', data = iris, hue = 'class', palette =
plt.title('Iris Dataset', fontsize = 14)
plt.legend(title = 'class', bbox_to_anchor=(1, 1))

plt.subplot(1,2,2)
sns.scatterplot(x = 'petallength', y = 'petalwidth', data = df, hue = 'cluster', palette =
plt.title('KMeans Model', fontsize = 14)
plt.legend(title = 'cluster', bbox_to_anchor=(1, 1))

plt.suptitle("Petalength vs Petalwidth", fontsize = 20)
plt.show()
```

executed in 1.16s, finished 18:14:57 2023-04-19



From both the plots, even though there were some mistakes, we can conclude that our model did a great job of creating the clusters for our data. We observe that cluster 1 represents setosa , cluster 0 represents versicolor and cluster 2 represents virginica .